

## concrete\_data

Shiva Sankar Modala

2023-03-01

```
#install.packages("tidyverse")

# readxl packages to load Excel data
#install.packages("readxl")
#install.packages("magrittr")
#install.packages("corrplot")

# Use the mgcv package to create a generalized additive model
#install.packages("mgcv")

# Visualize the regression using the visreg package,
#install.packages("visreg")

library(tidyverse)

## — Attaching core tidyverse packages — tidyverse
2.0.0 —
## ✓ dplyr      1.1.0      ✓ readr      2.1.4
## ✓ forcats   1.0.0      ✓ stringr    1.5.0
## ✓ ggplot2    3.4.1      ✓ tibble     3.1.8
## ✓ lubridate 1.9.2      ✓ tidyr      1.3.0
## ✓ purrr     1.0.1
## — Conflicts —
tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## ⓘ Use the http://conflicted.r-lib.org/conflicted-package to force
all conflicts to become errors

library(readxl)
library(magrittr)

##
## Attaching package: 'magrittr'
##
## The following object is masked from 'package:purrr':
##
##   set_names
##
## The following object is masked from 'package:tidyr':
##
##   extract
```

```

library(corrplot)

## corrplot 0.92 loaded

library(mgcv)

## Loading required package: nlme
##
## Attaching package: 'nlme'
##
## The following object is masked from 'package:dplyr':
##
##      collapse
##
## This is mgcv 1.8-41. For overview type 'help("mgcv-package")'.

library(visreg)

# Load the Concrete Compressive Strength sample dataset
concrete_data <- read_excel("C:/Users/shiva/OneDrive/Desktop/dpa
Assignments/Assignment3/Concrete_Data.xls")
summary(concrete_data)

## Cement (component 1)(kg in a m^3 mixture)
## Min.      :102.0
## 1st Qu.:192.4
## Median :272.9
## Mean    :281.2
## 3rd Qu.:350.0
## Max.    :540.0
## Blast Furnace Slag (component 2)(kg in a m^3 mixture)
## Min.      : 0.0
## 1st Qu.: 0.0
## Median : 22.0
## Mean     : 73.9
## 3rd Qu.:142.9
## Max.     :359.4
## Fly Ash (component 3)(kg in a m^3 mixture)
## Min.      : 0.00
## 1st Qu.: 0.00
## Median : 0.00
## Mean     : 54.19
## 3rd Qu.:118.27
## Max.     :200.10
## Water (component 4)(kg in a m^3 mixture)
## Min.      :121.8
## 1st Qu.:164.9
## Median :185.0
## Mean     :181.6
## 3rd Qu.:192.0
## Max.     :247.0

```

```


## Superplasticizer (component 5)(kg in a m^3 mixture)
## Min. : 0.000
## 1st Qu.: 0.000
## Median : 6.350
## Mean : 6.203
## 3rd Qu.:10.160
## Max. :32.200
## Coarse Aggregate (component 6)(kg in a m^3 mixture)
## Min. : 801.0
## 1st Qu.: 932.0
## Median : 968.0
## Mean : 972.9
## 3rd Qu.:1029.4
## Max. :1145.0
## Fine Aggregate (component 7)(kg in a m^3 mixture) Age (day)
## Min. :594.0 Min. : 1.00
## 1st Qu.:731.0 1st Qu.: 7.00
## Median :779.5 Median : 28.00
## Mean :773.6 Mean : 45.66
## 3rd Qu.:824.0 3rd Qu.: 56.00
## Max. :992.6 Max. :365.00
## Concrete compressive strength(MPa, megapascals)
## Min. : 2.332
## 1st Qu.:23.707
## Median :34.443
## Mean :35.818
## 3rd Qu.:46.136
## Max. :82.599

colnames(concrete_data) = c("cem", "bfs", "fa", "water", "sp", "cagg",
"fcagg", "age", "ccs")
column_names = c("cem", "bfs", "fa", "water", "sp", "cagg", "ccs")
concrete_data = concrete_data[column_names]
summary(concrete_data)

##          cem          bfs          fa          water
## Min. :102.0 Min. : 0.0 Min. : 0.00 Min. :121.8
## 1st Qu.:192.4 1st Qu.: 0.0 1st Qu.: 0.00 1st Qu.:164.9
## Median :272.9 Median : 22.0 Median : 0.00 Median :185.0
## Mean :281.2 Mean : 73.9 Mean : 54.19 Mean :181.6
## 3rd Qu.:350.0 3rd Qu.:142.9 3rd Qu.:118.27 3rd Qu.:192.0
## Max. :540.0 Max. :359.4 Max. :200.10 Max. :247.0
##          sp          cagg          ccs
## Min. : 0.000 Min. : 801.0 Min. : 2.332
## 1st Qu.: 0.000 1st Qu.: 932.0 1st Qu.:23.707
## Median : 6.350 Median : 968.0 Median :34.443
## Mean : 6.203 Mean : 972.9 Mean :35.818
## 3rd Qu.:10.160 3rd Qu.:1029.4 3rd Qu.:46.136
## Max. :32.200 Max. :1145.0 Max. :82.599

```

```
corrplot(cor(concrete_data), method = "number")
```



```
# gam function) to predict the Concrete Compressive Strength
dataModel1 <- gam(ccs ~ cem + bfs + fa + water + sp + cagg ,
data=concrete_data)
summary(dataModel1)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## ccs ~ cem + bfs + fa + water + sp + cagg
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.326997  10.510518   0.507 0.612387
## cem          0.108256   0.005214  20.761 < 2e-16 ***
## bfs          0.079357   0.006193  12.814 < 2e-16 ***
## fa           0.055928   0.009287   6.022 2.4e-09 ***
## water       -0.103871   0.027796  -3.737 0.000197 ***
## sp           0.356016   0.110251   3.229 0.001281 **
## cagg         0.008027   0.006272   1.280 0.200940
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
```

```
## R-sq.(adj) = 0.445   Deviance explained = 44.9%
## GCV = 155.83   Scale est. = 154.77   n = 1030

# compare the R2 value for a GAM with linear terms as well as smoothed terms
cat("The corrected R-squared + shows that a sizable portion of the variation
is present, and it appears that we have statistical effects for CEM and BFS
but not for CAGG.")

## The corrected R-squared + shows that a sizable portion of the variation is
present, and it appears that we have statistical effects for CEM and BFS but
not for CAGG.

# Use the s() function to apply smoothing using the default bs of tp).
dataModel2 <- gam(ccs ~ s(cem) + s(bfs) + s(fa) + s(water) + s(sp) + s(cagg)
, data=concrete_data)
summary(dataModel2)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## ccs ~ s(cem) + s(bfs) + s(fa) + s(water) + s(sp) + s(cagg)
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  35.8178      0.3566   100.4   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df      F  p-value
## s(cem)      4.464  5.513 69.530 < 2e-16 ***
## s(bfs)      2.088  2.578 48.091 < 2e-16 ***
## s(fa)       5.332  6.404  1.784  0.101
## s(water)    8.567  8.936 13.504 < 2e-16 ***
## s(sp)       7.133  8.143  5.498 1.22e-06 ***
## s(cagg)     1.000  1.000  0.018  0.892
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) = 0.531   Deviance explained = 54.4%
## GCV = 134.84   Scale est. = 130.96   n = 1030

cat("We should also remark that this model, with an adjusted R-squared
of.531, explains a large portion of the variance in CCS. In summary, it
appears that the CEM and CCS are connected.")

## We should also remark that this model, with an adjusted R-squared of.531,
explains a large portion of the variance in CCS. In summary, it appears that
the CEM and CCS are connected.
```

```

# showing the fit as a function of each predictor
dataModel1.sse <- sum(fitted(dataModel1)-concrete_data$ccs)^2
dataModel1.ssr <- sum(fitted(dataModel1) -mean(concrete_data$ccs))^2
dataModel1.sst = dataModel1.sse + dataModel1.ssr
Rsquared=1-(dataModel1.sse/dataModel1.sst)
cat(Rsquared)

## 0.4967177

dataModel2.sse <- sum(fitted(dataModel2)-concrete_data$ccs)^2
dataModel2.ssr <- sum(fitted(dataModel2) -mean(concrete_data$ccs))^2
dataModel2.sst = dataModel2.sse + dataModel2.ssr
Rsquared_sm=1-(dataModel2.sse/dataModel2.sst)
cat(Rsquared_sm)

## 0.5000744

anova(dataModel1, dataModel2, test="Chisq")


## Analysis of Deviance Table
##
## Model 1: ccs ~ cem + bfs + fa + water + sp + cagg
## Model 2: ccs ~ s(cem) + s(bfs) + s(fa) + s(water) + s(sp) + s(cagg)
##   Resid. Df Resid. Dev      Df Deviance   Pr(>Chi)
## 1    1023.00      158334
## 2     996.43      131019 26.574    27315 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

cat("Although we couldn't have known as much already, new statistical
evidence suggests that adding in the variables' nonlinear correlations
enhances the model.")


## Although we couldn't have known as much already, new statistical evidence
suggests that adding in the variables' nonlinear correlations enhances the
model.

visreg(dataModel1, 'cem')

```




```
visreg(dataModel2, 'cem')
```



```
cat("The end result, with all other model variables maintained constant,
    is a plot showing how the expected value of the CCS changes as a function
    of x (CEM).
    It contains the following information: (1) the expected value (blue
    line),
    (2) a confidence interval for the expected value (gray band), and
    (3) partial residuals (dark gray dots).")


## The end result, with all other model variables maintained constant,
## is a plot showing how the expected value of the CCS changes as a
## function of x (CEM).
## It contains the following information: (1) the expected value (blue
## line),
## (2) a confidence interval for the expected value (gray band), and
## (3) partial residuals (dark gray dots).

# Visualizing the feature with the function of their feature
visreg(dataModel1, 'bfs')
```




```
visreg(dataModel2, 'bfs')
```







```
visreg(dataModel1, 'fa')
```




```
visreg(dataModel2, 'fa')
```




```
visreg(dataModel1, 'water')
```




```
visreg(dataModel2, 'water')
```




```
visreg(dataModel1, 'sp')
```




```
visreg(dataModel2, 'sp')
```



```
visreg(dataModel1, 'cagg')
```



```
visreg(dataModel2, 'cagg')
```



```
cat("We can see from the CEM graph that the confidence interval has a higher  
value after adding the smoothing function than the model had without it.  
Using the smoothing function improves the confidence interval.")
```

```
## We can see from the CEM graph that the confidence interval has a higher  
value after adding the smoothing function than the model had without it.  
Using the smoothing function improves the confidence interval.
```